

# GRAFANA BEYLA (DAEMONSET)

## 1. Objective

Set up an observability stack for **Spring Boot applications** running on Kubernetes to:

- Collect **application metrics** (HTTP, gRPC, SQL, Redis, Messaging)
- Collect **process metrics** (CPU, memory, disk I/O, network)
- Expose metrics via **Prometheus**
- Visualize metrics in **Grafana**
- Ensure **high observability** and performance monitoring with minimal code changes

This setup is **non-intrusive**, using a **DaemonSet (Beyla)** to instrument applications automatically.

## 2. Components

Component	Purpose	Notes
Kubernetes Namespace	Isolate monitoring resources	namespace: beyla
ServiceAccount & RBAC	Grant Beyla permission to read Kubernetes metadata	Required for pod/service discovery
ConfigMap (Beyla)	Configuration file for instrumentation	/config/beyla-config.yml
DaemonSet (Beyla)	Runs the Beyla agent on every node	Instruments Java apps on port 8080
Prometheus	Scrapes metrics from Beyla	Exposed on port 9090
Grafana	Visualizes Prometheus metrics	Exposed on port 3000
Spring Boot Deployment	Application to monitor	Exposed on port 8080

## 3. Prerequisites

### Cluster Requirements

- Kubernetes v1.21+ (supports Lease API for leader election)
- RBAC enabled
- Ability to deploy DaemonSets
- NodePort or Ingress access for Prometheus/Grafana

- Minimum resources per node:
  - 500m CPU, 1GB RAM for monitoring stack (Prometheus + Beyla)

## Beyla Requirements

- Privileged mode enabled (SYS\_ADMIN, SYS\_PTRACE, BPF, etc.)
- Host PID access (hostPID: true) to instrument all processes
- Mount config files via ConfigMap

Proof:

- Kubernetes cluster - kubeadm (1 master + 2 workers) nodes should be ready.

```
pranavpp37@k8s-master:~/beyla$ kubectl get nodes
NAME           STATUS    ROLES          AGE   VERSION
k8s-master     Ready     control-plane  12d   v1.29.15
k8s-worker-1   Ready     <none>        12d   v1.29.15
k8s-worker-2   Ready     <none>        12d   v1.29.15
```

- pranavpp37/spring-boot-app: Docker image on all nodes

```
pranavpp37@k8s-master:~/beyla$ sudo docker images -a | grep pranavpp37/spring-boot-app
pranavpp37/spring-boot-app          v1              29ebbb35c055  8 days ago  204MB
pranavpp37@k8s-master:~/beyla$ |
```

- kubectl configured.

## IMPLEMENTATION

### Step – 1: Create Namespace

```
kubectl create namespace beyla
```

### Step-2: Deploy Beyla RBAC Permissions

```
01-beyla-namespace-rbac.yaml
---
apiVersion: v1
kind: Namespace
metadata:
  name: beyla
---
```

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: beyla
  namespace: beyla
---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: beyla
rules:
  - apiGroups: ["apps"]
    resources: ["replicasets"]
    verbs: ["list", "watch"]
  - apiGroups: [""]
    resources: ["pods", "services", "nodes"]
    verbs: ["list", "watch"]
---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: beyla
subjects:
  - kind: ServiceAccount
    name: beyla
    namespace: beyla
roleRef:
  kind: ClusterRole
  name: beyla
  apiGroup: rbac.authorization.k8s.io
```

Apply: kubectl apply -f 01-beyla-namespace-rbac.yaml

### Step-3: Deploy ConfigMap

```
02-beyla-configmap.yaml

---
apiVersion: v1
kind: ConfigMap
metadata:
  name: beyla-config
  namespace: beyla
data:
  beyla-config.yml: |
    # =====
    # BEYLA CONFIGURATION - PRODUCTION READY
    # =====

    # Enable Kubernetes metadata decoration
    # This adds pod names, namespaces, deployments to metrics
    attributes:
      kubernetes:
        enable: true

    # Configure routes to minimize cardinality
    # This groups similar URLs together (e.g., /user/123 -> /user/{id})
    routes:
      unmatched: heuristic

    # =====
    # SERVICE DISCOVERY - CRITICAL SECTION
    # =====

    # This tells Beyla which applications to instrument
    discovery:
```

```
services:

    # Find Java applications listening on port 8080
    # THIS WILL WORK - Your Spring Boot app uses port 8080
    - open_ports: 8080
    name: spring-boot-app

# =====
# OPENTELEMETRY METRICS EXPORT
# =====

otel_metrics_export:

    # If you have an OTEL collector, set endpoint here
    # Otherwise, leave as-is (uses env var)
    endpoint: ${OTEL_EXPORTER_OTLP_ENDPOINT}
    interval: 30s

# =====
# FEATURES - WHAT METRICS TO COLLECT
# =====

features:

    # WILL WORK - HTTP/gRPC RED metrics
    # Collects: request rate, errors, duration
    - application

    # WILL WORK - Process metrics
    # Collects: CPU, memory, disk I/O, network I/O
    - application_process

    # WILL WORK - Span-level metrics
    # Provides detailed per-request telemetry
    - application_span

# WILL WORK (if you have multiple services)
```

```
# Shows service-to-service communication  
  
# For single app, may have limited data  
- application_service_graph  
  
  
# WON'T GENERATE METRICS - Your app has no database  
  
# UNCOMMENT ONLY IF YOU ADD DATABASE TO YOUR APP  
  
# - sql  
  
  
# WON'T GENERATE METRICS - Your app doesn't use Redis  
  
# UNCOMMENT ONLY IF YOU ADD REDIS TO YOUR APP  
  
# - redis  
  
  
# WON'T GENERATE METRICS - Your app doesn't use MongoDB  
  
# UNCOMMENT ONLY IF YOU ADD MONGODB TO YOUR APP  
  
# - mongo  
  
  
# WON'T GENERATE METRICS - Your app doesn't use Kafka  
  
# UNCOMMENT ONLY IF YOU ADD KAFKA TO YOUR APP  
  
# - kafka  
  
  
# WON'T GENERATE METRICS - Your app doesn't use GPU  
  
# UNCOMMENT ONLY IF YOU HAVE GPU WORKLOADS  
  
# - gpu  
  
  
# =====  
# PROMETHEUS EXPORT - LOCAL SCRAPING  
# =====  
  
prometheus_export:  
  port: 9090  
  path: /metrics  
  
  
# Same features as OTEL export
```

```
features:
  - application          # HTTP metrics
  - application_process   # Process metrics
  - application_span       # Span metrics
  - application_service_graph # Service graph
  # - sql                  # Disabled - no database
  # - redis                # Disabled - no Redis
  # - mongo                # Disabled - no MongoDB
  # - kafka                # Disabled - no Kafka

# =====
# PROCESS METRICS CONFIGURATION
# =====

processes:
  enabled: true
  # These metrics will be collected:
  # process_cpu_time_seconds_total
  # process_cpu_utilization_ratio
  # process_memory_usage_bytes
  # process_memory_virtual_bytes
  # process_disk_io_bytes_total
  # process_network_io_bytes_total

# =====
# NETWORK METRICS - CURRENTLY DISABLED
# =====

# Network flow metrics are disabled to avoid conflicts
# Enable this ONLY if you want pod-to-pod network flows
# NOTE: Enabling this may put Beyla in network-only mode
# network:
#   enable: true
#
```

```
Apply: kubectl apply -f 02-beyla-configmap.yaml
```

#### Step-4: Deploy SpringBoot Java Application

```
03-java-app-deployment.yaml
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: spring-boot-app
  namespace: default
  labels:
    app: spring-boot-app
spec:
  replicas: 1
  selector:
    matchLabels:
      app: spring-boot-app
  template:
    metadata:
      labels:
        app: spring-boot-app
    spec:
      containers:
        - name: spring-boot
          image: pranavpp37/spring-boot-app:v1
          imagePullPolicy: IfNotPresent # Use local image, don't pull from registry
      ports:
        - containerPort: 8080
          name: http
      resources:
        requests:
```

```
        memory: "512Mi"
        cpu: "250m"
      limits:
        memory: "1Gi"
        cpu: "500m"
    ---
apiVersion: v1
kind: Service
metadata:
  name: spring-boot-svc
spec:
  type: NodePort
  selector:
    app: spring-boot-app
  ports:
    - port: 8080
      targetPort: 8080
```

Apply: kubectl apply -f 03-java-app-deployment.yaml

## Step-5: Deploy Beyla Daemonset

```
04-beyla-daemonset.yaml
---
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: beyla
  namespace: beyla
spec:
  selector:
    matchLabels:
```

```
instrumentation: beyla

template:

metadata:

labels:

instrumentation: beyla

spec:

serviceAccountName: beyla

hostPID: true # Required to see all processes on the node

containers:

- name: beyla

image: grafana/beyla:latest

securityContext:

privileged: true # Required for eBPF operations

capabilities:

add:

- SYS_ADMIN

- SYS_PTRACE

- SYS_RESOURCE

- NET_ADMIN

- NET_RAW

- BPF

- PERFMON

env:

# =====

# CRITICAL: Path to configuration file

# =====

- name: BEYLA_CONFIG_PATH

value: /config/beyla-config.yml

# =====

# Kubernetes metadata

# =====
```

```
- name: BEYLA_KUBE_METADATA_ENABLE
  value: "true"

# =====
# REMOVE ANY NETWORK-RELATED ENV VARS
# These were causing Beyla to run in network-only mode
# =====
# DO NOT SET: BEYLA_NETWORK_METRICS
# DO NOT SET: BEYLA_NETWORK_ENABLE

# =====
# OpenTelemetry endpoint (optional)
# Only set if you have an OTEL collector
# =====
# - name: OTEL_EXPORTER_OTLP_ENDPOINT
#   value: "http://your-otel-collector:4318"

# =====
# For Grafana Cloud (optional)
# =====
# - name: OTEL_EXPORTER_OTLP_ENDPOINT
#   value: "https://otlp-gateway-prod-us-central-0.grafana.net/otlp"
# - name: OTEL_EXPORTER_OTLP_HEADERS
#   value: "Authorization=Basic <base64-encoded-user:token>

# =====
# CRITICAL: Explicitly enable application metrics
# =====
- name: BEYLA_OTEL_METRICS_FEATURES
  value:
"application,application_process,application_span,application_service_graph"

# =====
```

```
# Log level for debugging
#
# -----
# - name: BEYLA_LOG_LEVEL
#   value: "INFO"

#
# -----
# Service discovery - CRITICAL
#
# -----
# - name: BEYLA_OPEN_PORT
#   value: "8080"

volumeMounts:
-
- name: config
  mountPath: /config

# Port for Prometheus metrics endpoint
ports:
-
- containerPort: 9090
  name: prometheus
  protocol: TCP

resources:
requests:
  memory: "200Mi"
  cpu: "100m"
limits:
  memory: "500Mi"
  cpu: "500m"

volumes:
-
- name: config
  configMap:
```

```
        name: beyla-config

        # Tolerate master node taints
        tolerations:
        - effect: NoSchedule
          key: node-role.kubernetes.io/control-plane
          operator: Exists
        ---

# Service to expose Prometheus metrics endpoint
apiVersion: v1
kind: Service
metadata:
  name: beyla-metrics
  namespace: beyla
  labels:
    instrumentation: beyla
spec:
  selector:
    instrumentation: beyla
  ports:
  - port: 9090
    targetPort: 9090
    name: prometheus
  type: ClusterIP

Apply: kubectl apply -f 04-beyla-daemonset.yaml
```

## Step-6: Deploy Prometheus

```
05-prometheus-setup.yaml

---
apiVersion: v1
kind: ConfigMap
metadata:
  name: prometheus-config
  namespace: beyla
data:
  prometheus.yml: |
    global:
      scrape_interval: 15s
      evaluation_interval: 15s

    scrape_configs:
      # Scrape Beyla metrics using service discovery
      - job_name: 'beyla'
        honor_labels: true
      kubernetes_sd_configs:
        - role: pod
          namespaces:
            names:
              - beyla
    relabel_configs:
      # Only scrape pods with instrumentation=beyla label
      - source_labels: [__meta_kubernetes_pod_label_instrumentation]
        action: keep
        regex: beyla
      # Use pod IP and port 9090
      - source_labels: [__meta_kubernetes_pod_ip]
        action: replace
        target_label: __address__
        replacement: ${1}:9090
```

```
# Add pod name as label
  - source_labels: [__meta_kubernetes_pod_name]
    action: replace
    target_label: pod

# Add node name as label
  - source_labels: [__meta_kubernetes_pod_node_name]
    action: replace
    target_label: node

# Alternative: Direct service endpoint scrape
  - job_name: 'beyla-service'
    honor_labels: true
    static_configs:
      - targets: ['beyla-metrics:9090']

---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: prometheus
  namespace: beyla
spec:
  replicas: 1
  selector:
    matchLabels:
      app: prometheus
  template:
    metadata:
      labels:
        app: prometheus
    spec:
      serviceAccountName: prometheus
      containers:
```

```
- name: prometheus
  image: prom/prometheus:latest
  args:
    - '--config.file=/etc/prometheus/prometheus.yml'
    - '--storage.tsdb.path=/prometheus'
    - '--web.console.libraries=/usr/share/prometheus/console_libraries'
    - '--web.console.templates=/usr/share/prometheus/consoles'
    - '--web.enable-lifecycle'
  ports:
    - containerPort: 9090
      name: web
  volumeMounts:
    - name: config
      mountPath: /etc/prometheus
    - name: storage
      mountPath: /prometheus
  resources:
    requests:
      memory: "512Mi"
      cpu: "250m"
    limits:
      memory: "2Gi"
      cpu: "1000m"
  volumes:
    - name: config
      configMap:
        name: prometheus-config
    - name: storage
      emptyDir: {}
---
apiVersion: v1
kind: ServiceAccount
```

```
metadata:
  name: prometheus
  namespace: beyla
---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: prometheus
rules:
  - apiGroups: [""]
    resources:
      - nodes
      - nodes/proxy
      - services
      - endpoints
      - pods
    verbs: ["get", "list", "watch"]
  - apiGroups: [""]
    resources:
      - configmaps
    verbs: ["get"]
---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: prometheus
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: prometheus
subjects:
  - kind: ServiceAccount
```

```
        name: prometheus
        namespace: beyla
    ---
apiVersion: v1
kind: Service
metadata:
    name: prometheus
    namespace: beyla
spec:
    selector:
        app: prometheus
    ports:
        - port: 9090
          targetPort: 9090
          nodePort: 30090
        name: web
    type: NodePort
```

Apply: kubectl apply -f 05-prometheus-setup.yaml

## Step-7: Deploy Grafana

```
06-grafana-setup.yaml
---
apiVersion: apps/v1
kind: Deployment
metadata:
    name: grafana
    namespace: beyla
spec:
    replicas: 1
    selector:
```

```
matchLabels:
  app: grafana

template:
  metadata:
    labels:
      app: grafana

spec:
  containers:
    - name: grafana
      image: grafana/grafana:latest
      ports:
        - containerPort: 3000
          name: web
      env:
        - name: GF_SECURITY_ADMIN_PASSWORD
          value: "admin" # Change this!
        - name: GF_INSTALL_PLUGINS
          value: "grafana-piechart-panel"
    volumeMounts:
      - name: storage
        mountPath: /var/lib/grafana
      - name: datasources
        mountPath: /etc/grafana/provisioning/datasources
  resources:
    requests:
      memory: "256Mi"
      cpu: "100m"
    limits:
      memory: "512Mi"
      cpu: "500m"
  volumes:
    - name: storage
```

```
    emptyDir: {}

    - name: datasources

      configMap:

        name: grafana-datasources

---


apiVersion: v1

kind: ConfigMap

metadata:

  name: grafana-datasources

  namespace: beyla

data:

  datasource.yml: |

    apiVersion: 1

    datasources:

      - name: Prometheus

        type: prometheus

        access: proxy

        url: http://prometheus:9090

        isDefault: true

        editable: true

---


apiVersion: v1

kind: Service

metadata:

  name: grafana

  namespace: beyla

spec:

  selector:

    app: grafana

  ports:

    - port: 3000

      targetPort: 3000
```

```

nodePort: 30000
type: NodePort

Apply: kubectl apply -f 06-grafana-setup.yaml

```

## VERIFICATION

) beyla namespace pods

```

pranavpp37@k8s-master:~/beyla$ kubectl get pods -n beyla -o wide
NAME          READY   STATUS    RESTARTS   AGE     IP           NODE      NOMINATED NODE   READINESS GATE
beyla-cmwm2   1/1    Running   1 (178m ago)  8d    192.168.230.21   k8s-worker-1   <none>        <none>
beyla-dfgv8   1/1    Running   1 (178m ago)  8d    192.168.235.204   k8s-master   <none>        <none>
beyla-vrv4f   1/1    Running   1 (178m ago)  8d    192.168.140.20    k8s-worker-2   <none>        <none>
grafana-dfd5c5c6b-crn8s  1/1    Running   1 (178m ago)  8d    192.168.230.18    k8s-worker-1   <none>        <none>
prometheus-5b5d89b996-g5z6k  1/1    Running   1 (178m ago)  8d    192.168.230.19    k8s-worker-1   <none>        <none>
pranavpp37@k8s-master:~/beyla$ 

```

) default (java-app pods)

```

spring-boot-app-5b66f9659-ljmvf  1/1    Running   1 (178m ago)  8d    192.168.230.20   k8s-worker-1   <none>        <none>

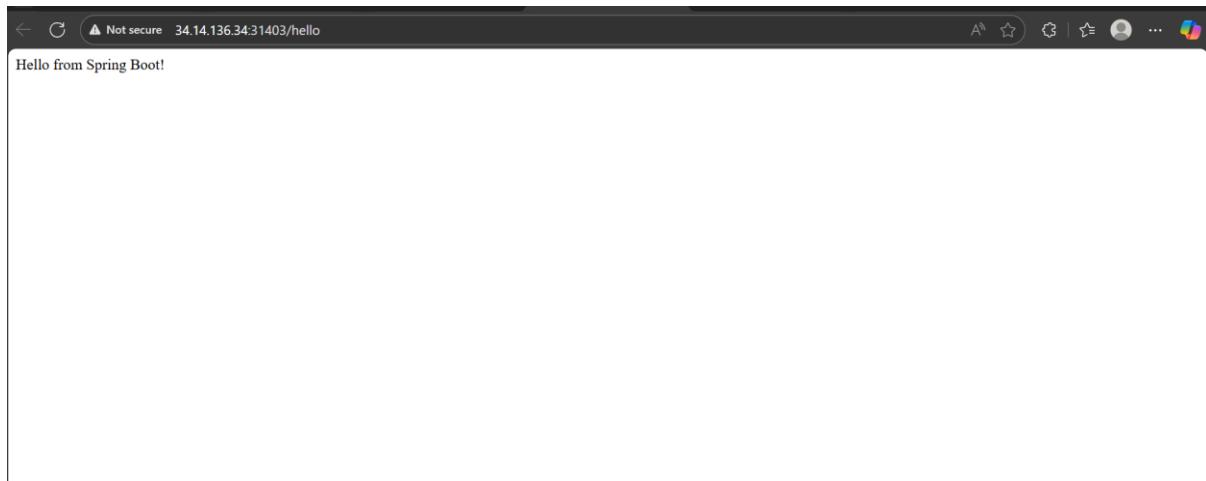
```

## ENDPOINTS

Service / App	Namespace	NodePort / Port	Access URL	Notes
Spring Boot App	default	8080 (ClusterIP)	http://<Node-IP>:8080/hello via port-forward	Use port-forward since NodePort not defined
Prometheus	beyla	30090	http://34.14.136.34:30090	Scrapes Beyla metrics at /metrics
Grafana	beyla	30000	http://34.14.136.34:30000	Default login: admin/admin
Beyla metrics	beyla	9090 (pod)	http://<pod-ip>:9090/metrics	Access via port-forward

## SCREENSHOTS

) Java App



## ) Prometheus

The screenshot shows the Prometheus web interface with three main sections:

- Targets:** A table listing three targets under the job "beyla". Each target has multiple endpoints (metrics) listed with their last scrape time and state (UP).

Endpoint	Labels	Last scrape	State
http://192.168.235.204:9090/metrics	instance="192.168.235.204:9090" job="beyla" node="k8s-master" pod="beyla-dfgv8"	10.77s ago	1ms UP
http://192.168.140.20:9090/metrics	instance="192.168.140.20:9090" job="beyla" node="k8s-worker-2" pod="beyla-vrv4l"	12.974s ago	2ms UP
http://192.168.230.21:9090/metrics	instance="192.168.230.21:9090" job="beyla" node="k8s-worker-1" pod="beyla-cmwm2"	10.012s ago	4ms UP
- beyla-service:** A table listing one target under the job "beyla-service".

Endpoint	Labels	Last scrape	State
http://beyla-metrics:9090/metrics	instance="beyla-metrics:9090" job="beyla-service"	5.334s ago	1ms UP
- Metrics Explorer:** A search interface for exploring available metrics. It includes a search bar, a table of metrics with their descriptions, and a "Execute" button.

Metric	Description
http_server_request_duration_seconds_sum	duration of HTTP service calls from the server side, in seconds
http_server_response_body_size_bytes_bucket	size, in bytes, of the HTTP response body as received at the server side
http_server_response_body_size_bytes_count	size, in bytes, of the HTTP response body as received at the server side
http_server_response_body_size_bytes_sum	size, in bytes, of the HTTP response body as received at the server side
process_cpu_time_seconds_total	Total CPU seconds broken down by different states
process_cpu_utilization_ratio	Difference in process.cpu.time since the last measurement, divided by the elapsed time and number of CPUs available to the process
process_disk_io_bytes_total	Disk bytes transferred
process_memory_usage_bytes	The amount of physical memory in use
process_memory_virtual_bytes	The amount of committed virtual memory
process_network_io_bytes_total	Network bytes transferred
promhttp_metric_handler_errors_total	Total number of internal errors encountered by the promhttp metric handler.
scrape_duration_seconds	unknown
scrape_samples_post_metric_relabeling	unknown

The screenshot shows the Grafana web interface with a query editor. The left sidebar shows a list of targets and a table of metrics. The main area shows a table of metrics with their descriptions and a search bar.

Metric	Description
http_server_request_duration_seconds_sum	duration of HTTP service calls from the server side, in seconds
http_server_response_body_size_bytes_bucket	size, in bytes, of the HTTP response body as received at the server side
http_server_response_body_size_bytes_count	size, in bytes, of the HTTP response body as received at the server side
http_server_response_body_size_bytes_sum	size, in bytes, of the HTTP response body as received at the server side
process_cpu_time_seconds_total	Total CPU seconds broken down by different states
process_cpu_utilization_ratio	Difference in process.cpu.time since the last measurement, divided by the elapsed time and number of CPUs available to the process
process_disk_io_bytes_total	Disk bytes transferred
process_memory_usage_bytes	The amount of physical memory in use
process_memory_virtual_bytes	The amount of committed virtual memory
process_network_io_bytes_total	Network bytes transferred
promhttp_metric_handler_errors_total	Total number of internal errors encountered by the promhttp metric handler.
scrape_duration_seconds	unknown
scrape_samples_post_metric_relabeling	unknown

## ) Grafana

