

Kubernetes Lease – Kubernetes Leader-Elector Sidecar Setup

What is Kubernetes Lease?

In Kubernetes, a Lease is a special resource (kind: Lease) in the **coordination.k8s.io/v1** API group that is used to implement leader election or distributed coordination among multiple pods.

A **Lease** represents a lightweight lock that a pod can acquire for a limited time to become the “leader.” It contains metadata such as the current holder, acquisition time, and lease duration. Other pods can attempt to acquire the lease when it expires, ensuring that at most one pod is the leader at a time.

Key points:

- i) Stored in the Kubernetes API server (backed by etcd).
- ii) Usually used for **leader election** by controllers or custom sidecars.
- iii) Helps prevent multiple replicas from performing the same operation simultaneously.
- iv) The lease has fields like:
 - holderIdentity - the current leader's identifier
 - leaseDurationSeconds - how long the lease is valid
 - acquireTime - when the lease was acquired
 - renewTime - last time the leader renewed the lease

NOTE: Kubernetes Lease = a lightweight API-based lock for leader election among pods.

APPROACHES

- i) **Use a Kubernetes client library (Java, Go, Python, etc.) in your app.**

Example client libraries:

- Java: Fabric8 Kubernetes Client → LeaderElector with Lease
- Go: client-go → leaderelection package with LeaseLock
- Python: kubernetes library → custom Lease acquisition.

For Further Reference: <https://www.linkedin.com/pulse/kubernetes-leases-leader-election-ensuring-high-suyash-gogte-tybbf>

ii) Kubernetes Lease API Sidecar Leader Election

Pre-Requisites

- i) A running **Kubernetes cluster** (kubeadm-based is fine).
- ii) kubectl configured and pointing to your cluster.
- iii) At least v1.13+ Kubernetes (because leader election via Endpoints is supported from early versions).
- iv) 3 replicas of your application (or more) to test leader election.
- v) Main image: (e.g., nginx:1.25 as placeholder).
- vi) Leader-elector sidecar image: ‘gcr.io/google_containers/leader-elector:0.4’.

Implementation

rbac.yaml

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: leader-elector-sa
  namespace: default
---
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  name: leader-elector-role
  namespace: default
rules:
# Endpoints-based leader election (required for gcr.io/google_containers/leader-elector:0.4)
- apiGroups: [""]
  resources: ["endpoints"]
  verbs: ["get", "create", "update", "patch"]
# Optional: pods read (some modes read pod identity)
- apiGroups: [""]
  resources: ["pods"]
  verbs: ["get", "list", "watch"]
```

```
# Optional: future-proof if you ever switch to Lease
- apiGroups: ["coordination.k8s.io"]
  resources: ["leases"]
  verbs: ["get", "create", "update", "patch"]

---
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: leader-elector-binding
  namespace: default
subjects:
- kind: ServiceAccount
  name: leader-elector-sa
  namespace: default
roleRef:
  kind: Role
  name: leader-elector-role
  apiGroup: rbac.authorization.k8s.io
```

Apply: kubectl apply -f rbac.yaml

deployment.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-app
  namespace: default
spec:
  replicas: 3
  selector:
    matchLabels:
      app: my-app
```

```

template:

metadata:
  labels:
    app: my-app

spec:
  serviceAccountName: leader-elector-sa

  containers:
    - name: main-app
      image: nginx:1.25

    - name: leader-elector
      image: gcr.io/google_containers/leader-elector:0.4
      args:
        - "--election=leader-election"    # Endpoints object name in 'default' ns
        - "--http=0.0.0.0:4040"          # sidecar status endpoint
        - "--ttl=30s"                   # renew period / ttl

  ports:
    - name: http
      containerPort: 4040

```

Apply: kubectl apply -f deployment.yaml

VERIFICATION

i) Command: kubectl get po -A

default	my-app-789986b684-8zlbx	2/2	Running	0	39m
default	my-app-789986b684-9hck7	2/2	Running	0	38m
default	my-app-789986b684-p8nfx	2/2	Running	0	39m

ii) Leader info via Endpoints annotation

```
Command: kubectl get endpoints leader-election -n default -o yaml
```

```
pranavpp37@k8s-master:~$ kubectl get endpoints leader-election -n default -o yaml
apiVersion: v1
kind: Endpoints
metadata:
  annotations:
    control-plane.alpha.kubernetes.io/leader: '{"holderIdentity":"my-app-789986b684-8zlbx","leaseDurationSeconds":30,"acquireTime":"2025-12-29T11:10:58Z","renewTime":"2025-12-29T11:50:31Z","LeaderTransitions":0}'
    creationTimestamp: "2025-12-29T11:08:02Z"
  name: leader-election
  namespace: default
  resourceVersion: "791440"
  uid: ac8ad671-55ca-4f88-b155-e8cca5ede95d
```

iii) Extract leader pod name

```
Command: kubectl get endpoints leader-election -o
jsonpath='{.metadata.annotations.control-plane\\.alpha\\.kubernetes\\.io/leader}' \
| sed -E 's/.+"holderIdentity":(["]+).*/\1/' && echo
```

```
pranavpp37@k8s-master:~$ kubectl get endpoints leader-election -o jsonpath='{.metadata.annotations.control-plane\\.alpha\\.kubernetes\\.io/leader}' \
| sed -E 's/.+"holderIdentity":(["]+).*/\1/' && echo
my-app-789986b684-8zlbx
```

iv) Check logs of Leader

```
Command: kubectl logs my-app-789986b684-8zlbx
```

```
pranavpp37@k8s-master:~$ kubectl logs my-app-789986b684-8zlbx
Defaulted container "main-app" out of: main-app, leader-elector
/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
/docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
/docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
10-listen-on-ipv6-by-default.sh: info: Getting the checksum of /etc/nginx/conf.d/default.conf
10-listen-on-ipv6-by-default.sh: info: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
/docker-entrypoint.sh: Sourcing /docker-entrypoint.d/15-local-resolvers.envsh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/30-tune-worker-processes.sh
/docker-entrypoint.sh: Configuration complete; ready for start up
2025/12/29 11:09:53 [notice] 1#1: using the "epoll" event method
2025/12/29 11:09:53 [notice] 1#1: nginx/1.25.5
2025/12/29 11:09:53 [notice] 1#1: built by gcc 12.2.0 (Debian 12.2.0-14)
2025/12/29 11:09:53 [notice] 1#1: OS: Linux 6.8.0-1045-gcp
2025/12/29 11:09:53 [notice] 1#1: getrlimit(RLIMIT_NOFILE): 1048576:1048576
2025/12/29 11:09:53 [notice] 1#1: start worker processes
2025/12/29 11:09:53 [notice] 1#1: start worker process 29
2025/12/29 11:09:53 [notice] 1#1: start worker process 30
```

Command: kubectl logs my-app-789986b684-8zlbx -c leader-elector

```
pranavpp370k8s-master:~$ kubectl logs my-app-789986b684-8zlbx -c leader-elector
my-app-ddbc6cb96-wjnlf is the leader
I1229 11:09:53.185992      8 leaderelection.go:259] lock is held by my-app-ddbc6cb96-wjnlf and has not yet expired
I1229 11:10:06.132517      8 leaderelection.go:259] lock is held by my-app-ddbc6cb96-wjnlf and has not yet expired
I1229 11:10:22.101432      8 leaderelection.go:259] lock is held by my-app-ddbc6cb96-wjnlf and has not yet expired
I1229 11:10:35.587057      8 leaderelection.go:259] lock is held by my-app-ddbc6cb96-wjnlf and has not yet expired
I1229 11:10:47.031278      8 leaderelection.go:259] lock is held by my-app-ddbc6cb96-wjnlf and has not yet expired
my-app-789986b684-8zlbx is the leader
E1229 11:10:58.364566      8 event.go:247] Could not construct reference to: '&api.Endpoints{TypeMeta:unversioned.TypeMeta{Kind:"", APIVersion:""}, ObjectMeta:api.ObjectMeta{Name:"leader-election", GenerateName:"", Namespace:"default", SelfLink:"", UID:"ac8ad671-55ca-4f88-b155-e8cca5ede95d", ResourceVersion:"787275", Generation:0, CreationTimestamp:unversioned.Time{Time:time.Time{sec:63902603282, nsec:0, loc:(*time.Location)(0x19072e0)}}, DeletionTimestamp:(*unversioned.Time)(nil), DeletionGracePeriodSeconds:(*int64)(nil), Labels:map[string]string(nil), Annotations:map[string]string{"control-plane.alpha.kubernetes.io/Leader":{"\"holderIdentity\":\"my-app-ddbc6cb96-wjnlf\"}}, "leaseDurationSeconds":30, "acquireTime\":\"2025-12-29T11:08:02Z\", "renewTime\":\"2025-12-29T11:09:47Z\", "LeaderTransitions\":0}}}, Subsets:[]api.EndpointSubset(nil)}' due to: 'selfLink was empty, can't make reference'. Will not report event: 'Normal' '%v became leader' 'my-app-789986b684-8zlbx'
I1229 11:10:58.364782      8 leaderelection.go:180] sucessfully acquired lease default/leader-election
pranavpp370k8s-master:~$ |
```

v) Check logs of other pods

Command: kubectl logs my-app-789986b684-9hck7 -c leader-elector

```
I1229 11:53:26.035041      8 leaderelection.go:259] lock is held by my-app-789986b684-8zlbx and has not yet expired
I1229 11:53:36.413893      8 leaderelection.go:259] lock is held by my-app-789986b684-8zlbx and has not yet expired
I1229 11:53:48.413910      8 leaderelection.go:259] lock is held by my-app-789986b684-8zlbx and has not yet expired
I1229 11:53:59.522062      8 leaderelection.go:259] lock is held by my-app-789986b684-8zlbx and has not yet expired
I1229 11:54:07.204637      8 leaderelection.go:259] lock is held by my-app-789986b684-8zlbx and has not yet expired
I1229 11:54:20.514276      8 leaderelection.go:259] lock is held by my-app-789986b684-8zlbx and has not yet expired
I1229 11:54:31.876585      8 leaderelection.go:259] lock is held by my-app-789986b684-8zlbx and has not yet expired
I1229 11:54:42.437327      8 leaderelection.go:259] lock is held by my-app-789986b684-8zlbx and has not yet expired
I1229 11:54:57.928738      8 leaderelection.go:259] lock is held by my-app-789986b684-8zlbx and has not yet expired
I1229 11:55:07.559915      8 leaderelection.go:259] lock is held by my-app-789986b684-8zlbx and has not yet expired
I1229 11:55:21.949250      8 leaderelection.go:259] lock is held by my-app-789986b684-8zlbx and has not yet expired
I1229 11:55:29.774831      8 leaderelection.go:259] lock is held by my-app-789986b684-8zlbx and has not yet expired
I1229 11:55:43.827615      8 leaderelection.go:259] lock is held by my-app-789986b684-8zlbx and has not yet expired
I1229 11:55:56.966115      8 leaderelection.go:259] lock is held by my-app-789986b684-8zlbx and has not yet expired
I1229 11:56:09.088473      8 leaderelection.go:259] lock is held by my-app-789986b684-8zlbx and has not yet expired
I1229 11:56:17.244438      8 leaderelection.go:259] lock is held by my-app-789986b684-8zlbx and has not yet expired
I1229 11:56:31.267448      8 leaderelection.go:259] lock is held by my-app-789986b684-8zlbx and has not yet expired
I1229 11:56:46.690334      8 leaderelection.go:259] lock is held by my-app-789986b684-8zlbx and has not yet expired
```